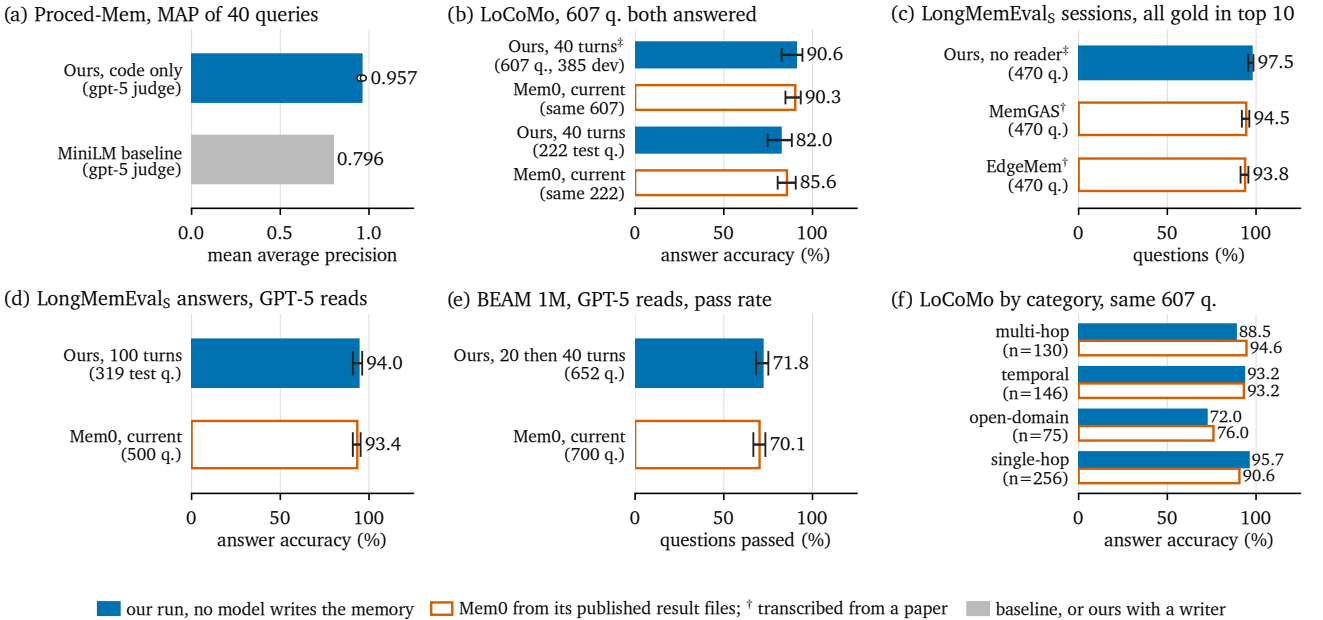
(a) Proced-Mem, MAP of 40 queries
Ours, code only (gpt-5 judge)
0.957
MiniLM baseline (gpt-5 judge)
0.796
0.0
0.5
1.0
mean average precision
(b) LoCoMo, 607 q. both answered
Ours, 40 turns‡ (607 q., 385 dev)
90.6
Mem0, current (same 607)
90.3
Ours, 40 turns (222 test q.)
82.0
Mem0, current (same 222)
85.6
0
50
100
answer accuracy (%)
(c) LongMemEval$_S$ sessions, all gold in top 10
Ours, no reader‡ (470 q.)
97.5
MemGAS† (470 q.)
94.5
EdgeMem† (470 q.)
93.8
0
50
100
questions (%)
(d) LongMemEval$_S$ answers, GPT-5 reads
Ours, 100 turns (319 test q.)
94.0
Mem0, current (500 q.)
93.4
0
50
100
answer accuracy (%)
(e) BEAM 1M, GPT-5 reads, pass rate
Ours, 20 then 40 turns (652 q.)
71.8
Mem0, current (700 q.)
70.1
0
50
100
questions passed (%)
(f) LoCoMo by category, same 607 q.
multi-hop (n=130)
88.5
94.6
temporal (n=146)
93.2
93.2
open-domain (n=75)
72.0
76.0
single-hop (n=256)
95.7
90.6
0
50
100
answer accuracy (%)
our run, no model writes the memory
Mem0 from its published result files; † transcribed from a paper
baseline, or ours with a writer